

TSENAT: Tsallis Entropy Analysis Toolbox

Cristóbal Gallardo gallardoalba@pm.me

2026-02-10

Overview

This vignette shows how to compute and apply Tsallis entropy to transcript-level expression data with TSENAT. It focuses on three practical goals: (1) computing per-sample, per-gene diversity measures across a range of sensitivity parameters q ; (2) comparing those measures between sample groups; and (3) visualizing and inspecting transcript-level patterns that explain observed differences.

Motivation

Common RNA-seq tools, such as DESEQ2 or SALMON, focus on changes in total gene abundance. While powerful, that view can obscure an important biological phenomenon: genes often change *which* isoforms they express without large changes in overall expression. Tsallis entropy gives you a single, coherent framework to describe those changes. By tuning the parameter q you slide a lens across abundance scales—zooming in on rare variants or zooming out to the dominant isoforms—so you can capture scale-dependent isoform switching that ordinary summaries overlook.

High-level workflow

1. Preprocess counts and filter low-abundance transcripts.
2. Compute relative transcript proportions within each gene and evaluate Tsallis entropy across one or more q values.
3. Summarize per-gene values across samples (mean/median) and test for differences between groups (Wilcoxon or permutations).
4. Visualize q -curves and inspect transcript-level counts for top genes.

What is Tsallis entropy?

Tsallis entropy is a one-parameter family of diversity measures that generalizes Shannon entropy via a sensitivity parameter q . For a discrete probability vector $p = (p_1, \dots, p_n)$ it is defined as

$$S_q(p) = \frac{1 - \sum_{i=1}^n p_i^q}{q - 1}.$$

In information theory terms, entropy measures the uncertainty when drawing a single transcript from the gene's isoform frequency distribution: higher entropy means the choice is less predictable (many similarly abundant isoforms), while lower entropy means the distribution is dominated by one or a few transcripts.

Tsallis entropy allows, through the q parameter, shifts emphasis between rare and abundant isoforms: small values of q make the measure sensitive to low-abundance variants, while larger q focus on the dominant

species. In practice, computing S_q across a short grid of q values (a “ q -curve”) and combining this with resampling or paired designs yields a concise, interpretable summary of isoform heterogeneity.

Essential limiting cases:

- Limit $q \rightarrow 1$: Shannon entropy, $\lim_{q \rightarrow 1} S_q(p) = -\sum_i p_i \log p_i$.
- $q = 0$: richness-like (number of expressed isoforms minus one).
- $q = 2$: Gini–Simpson / collision index, $S_2 = 1 - \sum_i p_i^2$.
- Uniform maximum: for m expressed isoforms, $S_{q,\max}(m) = \frac{1-m^{1-q}}{q-1}$ and normalized $\tilde{S}_q = S_q/S_{q,\max}(m)$.

Example dataset

An example dataset is included for demonstration.

```
# Load packages
suppressPackageStartupMessages({
  library(TSENAT)
  library(ggplot2)
  library(SummarizedExperiment)
})
```

Load data and metadata

Now we will load the example dataset and associated metadata:

```
# Load required files
coldata_tsv <- system.file("extdata", "coldata.tsv", package = "TSENAT")
tx2gene_tsv <- system.file("extdata", "tx2gene.tsv", package = "TSENAT")
data("tcga_brca_luma_dataset", package = "TSENAT", envir = globalenv())

# Extract gene names and read count data
genes <- tcga_brca_luma_dataset[, 1]
readcounts <- tcga_brca_luma_dataset[, -1]

# Read sample metadata
coldata_df <- read.table(coldata_tsv,
  header = TRUE,
  sep = "\t",
  stringsAsFactors = FALSE
)
```

The next step is to inspect the loaded data to ensure it looks correct.

```
# Check gene names
head(genes)
#> [1] "MXRA8"    "MXRA8"    "MXRA8"    "MXRA8"    "MXRA8"    "C1orf86"

# Check read count dataset
dim(readcounts)
#> [1] 1100   40
head(readcounts[1:4, 1:3])
```

```

#>   TCGA-A7-AOCH_N TCGA-A7-AOCH_T TCGA-A7-AOD9_N
#> 1      2858.04      743.56      812.59
#> 2      127.82       21.28      50.87
#> 3      370.22       94.38      368.76
#> 4     7472.00      3564.87     7647.76

# Check the metadata
head(coldata_df)
#>           Sample Condition
#> 1 TCGA-A7-AOCH_N    Normal
#> 2 TCGA-A7-AOD9_N    Normal
#> 3 TCGA-A7-AODB_N    Normal
#> 4 TCGA-A7-AODC_N    Normal
#> 5 TCGA-A7-A13G_N    Normal
#> 6 TCGA-AC-A2FB_N    Normal

```

Data preprocessing and filtering

If the transcript IDs in `readcounts` are not already set as rownames, we need to map them. This step could be skipped if the dataset already has transcript IDs as rownames, but in this case we need to assign them.

```

# Assign transcript IDs as rownames of `readcounts`
readcounts <- map_tx_to_readcounts(readcounts, tx2gene_tsv)

```

Next, we will create a `SummarizedExperiment` data container, a standard Bioconductor class for storing high-throughput assay data along with associated metadata. You can read more about it in the Bioconductor documentation.

```

## Build a `SummarizedExperiment` from readcounts + tx2gene mapping
se <- build_se(tx2gene_tsv, readcounts, genes)

```

Before computing the Tsallis diversity, you might want to filter out genes with a low overall expression or limit the analysis to transcripts with a sufficient minimum expression level. Expression estimates of transcript isoforms with zero or low expression might be highly variable. For more details on the effect of transcript isoform prefiltering on differential transcript usage, see this paper.

```

## Filter lowly-expressed transcripts and report counts
# Keep transcripts with >5 reads in more than 5 samples using `filter_se()`
se <- filter_se(se, min_count = 5, min_samples = 5)
#> Transcripts: before = 1100, after = 848

```

We filtered lowly-expressed transcripts to reduce noise and improve the stability of diversity estimates.

Compute normalized Tsallis entropy.

Now we will compute Tsallis entropy for a single q , and inspect the resulting assay.

Note: For single-condition studies (no comparison groups), Tsallis entropy can still be computed and used for descriptive analyses and quality control; formal group-comparison tests are not applicable in this setting.

```
# compute Tsallis entropy for a single q value (normalized)
q <- 0.1
ts_se <- calculate_diversity(se, q = q, norm = TRUE)
head(assay(ts_se)[1:4, 1:3])
#>      TCGA-A7-AOCH_N TCGA-A7-AOCH_T TCGA-A7-AOD9_N
#> MXRA8          0.8616275   0.6609111   0.8156972
#> C1orf86        0.0000000   0.0000000   0.0000000
#> PDPN           0.2657489   0.2588286   0.0000000
#> ALDH4A1         0.8111864   0.4667361   0.8621495
```

An additional important technical factor is the data normalization, which makes values comparable across genes with different numbers of isoforms.

Now we should map the sample metadata into the `SummarizedExperiment` object (`ts_se`), so plotting functions and downstream helpers can rely on it.

```
ts_se <- map_metadata(ts_se, coldata_df, paired = TRUE)
```

Differential analysis

This section summarizes practical guidance for hypothesis testing and reporting when comparing diversity measures across sample groups.

The package integrate diverse tests for comparing per-gene diversity summaries between two groups. Two common choices are Wilcoxon-based tests and label-shuffle tests.

Here we will use the **Wilcoxon test**; it is the appropriate choice for speed when sample sizes support asymptotic approximations, and it is robust to non-normality. The **label-shuffling test** is a non-parametric alternative that does not rely on distributional assumptions and can be more accurate for small sample sizes, but it is computationally intensive.

For the analyses shown in this vignette we recommend using the **median** in the **Wilcoxon** rank-sum test for comparisons. The reason is that Tsallis entropy values are often skewed and sensitive to outliers, and the median provides in this condition a more robust summary of central tendency than the mean.

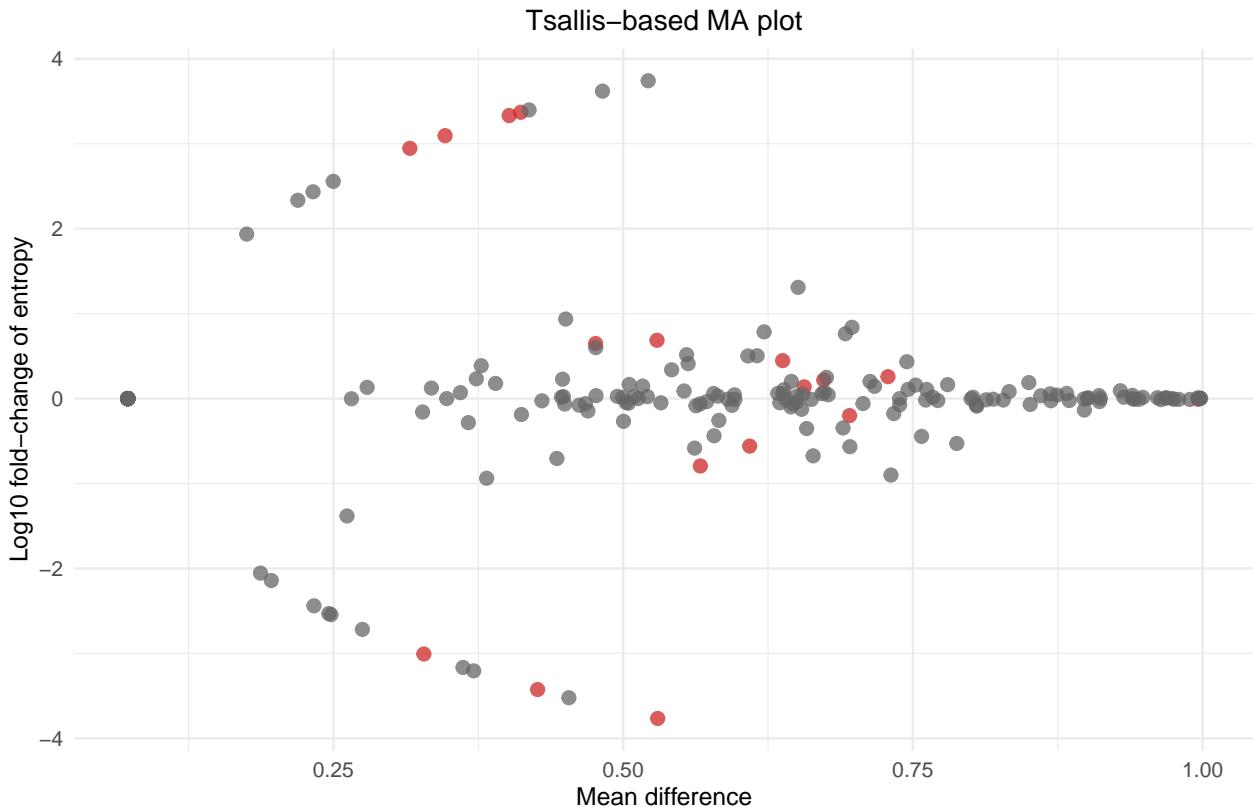
```
# Compute differences between groups using the Wilcoxon test
# with median aggregation
res <- calculate_difference(
  ts_se,
  control = "Normal",
  method = "median",
  test = "wilcoxon",
  paired = TRUE
)
# sort results by adjusted p-value
res <- res[order(res$adjusted_p_values), ]
head(res)
#>      genes Tumor_median Normal_median median_difference log2_fold_change
#> 6    C1orf213  0.07256896  0.98702527   -0.91445631   -3.765662496
#> 17   S1PR1    0.99315946  0.99930867   -0.00614921   -0.008904999
#> 88    MBD2    0.55929534  0.07256896    0.48672638    2.946185792
#> 99    OSR1    0.79339202  0.66364511    0.12974692    0.257621943
#> 102   GFPT1    0.07256896  0.58345094   -0.51088198   -3.007186825
#> 128   KIF9    0.58141625  0.37105427    0.21036198    0.647941169
```

```
#>      raw_p_values adjusted_p_values
#> 6  0.0002357454  0.02721101
#> 17 0.0007285418  0.02721101
#> 88 0.0006535852  0.02721101
#> 99 0.0005801974  0.02721101
#> 102 0.0008919014 0.02721101
#> 128 0.0006356242  0.02721101
```

After the statistical test, in order to visualize the results and inspect the most significant genes, we can generate two kind of diagnostic plots: MA and volcano plots.

First we will plot a **Tsallis-based MA**, which uses the Tsallis entropy values directly to compute fold changes and mean expression. This plot provides a direct view of how diversity differences relate to overall expression levels.

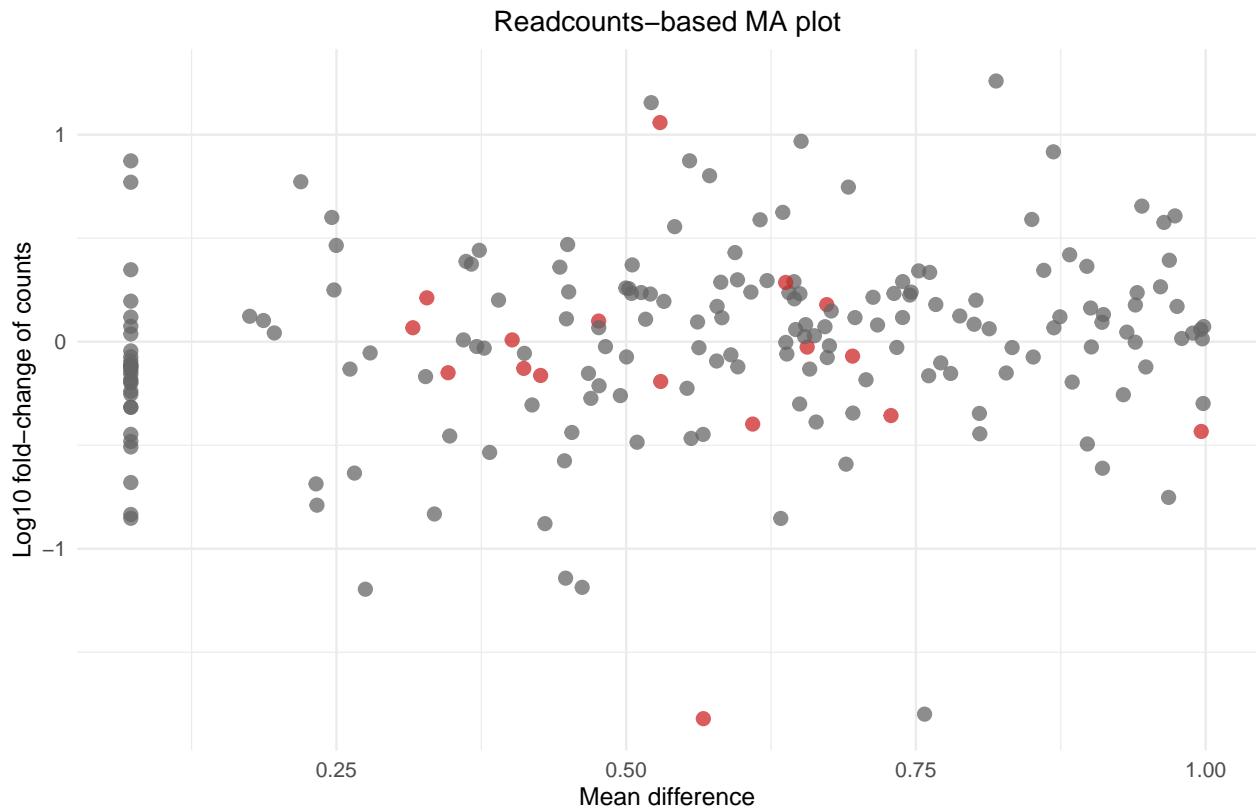
```
# MA plot (Tsallis-based)
p_ma <- plot_ma_tsallis(res)
print(p_ma)
```



Second, we will plot an **expression/readcount-based MA**, which uses the original read counts to compute fold changes and mean expression. This plot helps to contextualize diversity differences in terms of overall gene expression levels.

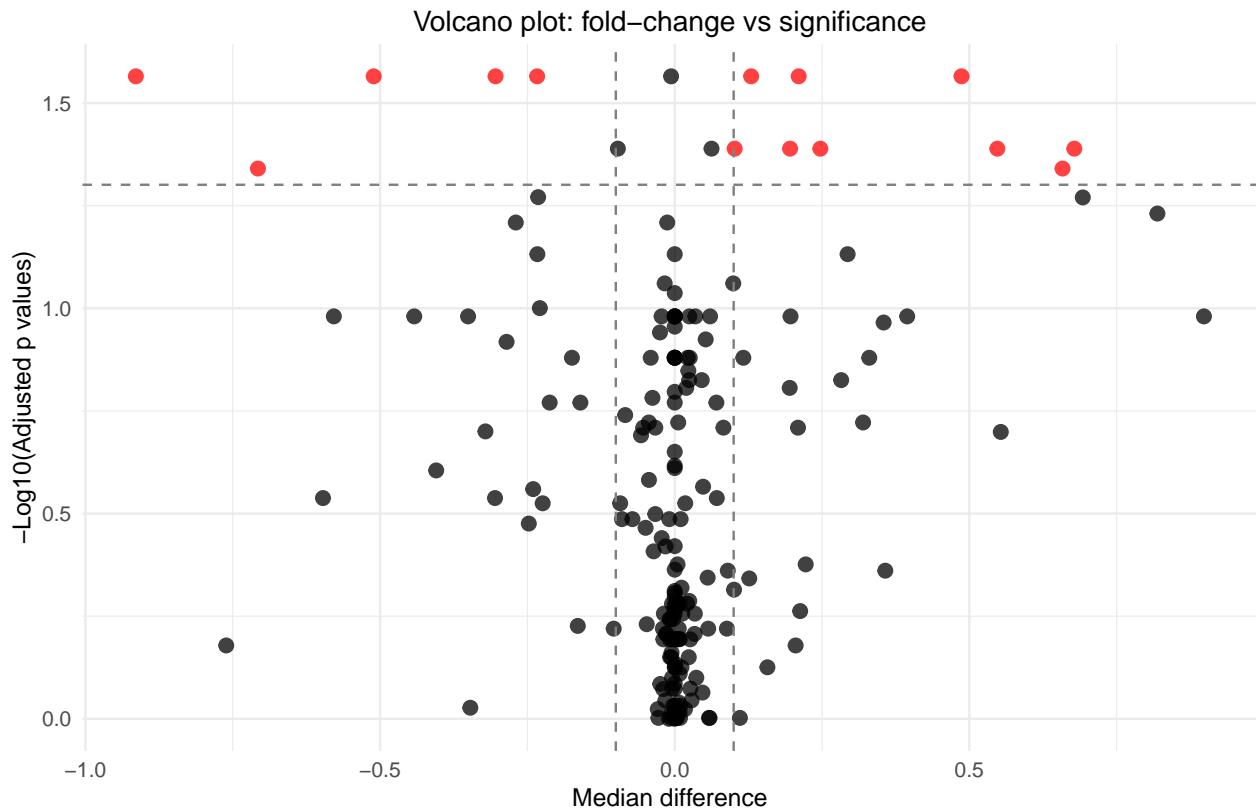
```
# MA plot (expression/readcount-based)
p_ma_read <- plot_ma_expression(
  res,
  se = ts_se,
  control = "Normal")
```

```
)  
print(p_ma_read)
```



On the other hands, the **volcano plot** displays the relationship between the magnitude of diversity differences (mean difference) and their statistical significance (adjusted p-value). This plot helps to identify genes with both large effect sizes and strong statistical support.

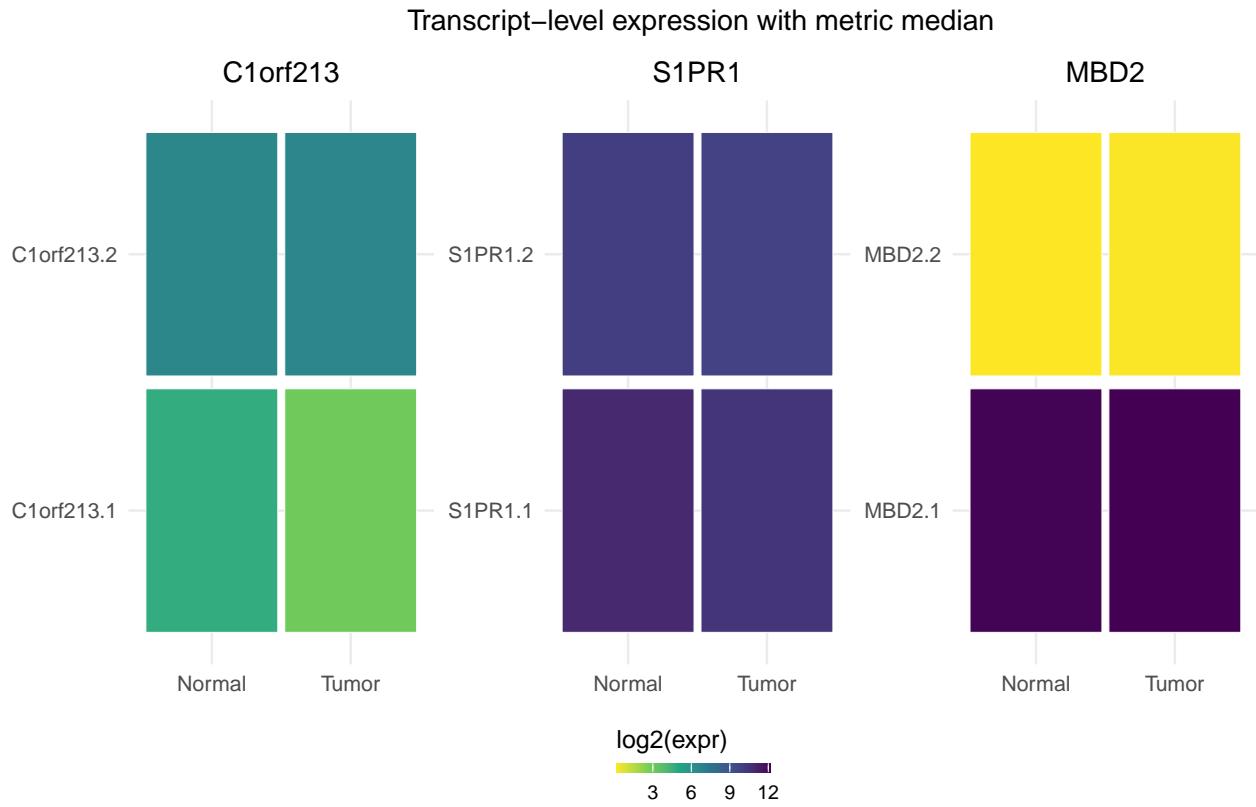
```
# Volcano plot: mean difference vs -log10(adjusted p-value)  
p_volcano <- plot_volcano(res)  
print(p_volcano)
```



Plot top 3 genes from the single-q differential analysis

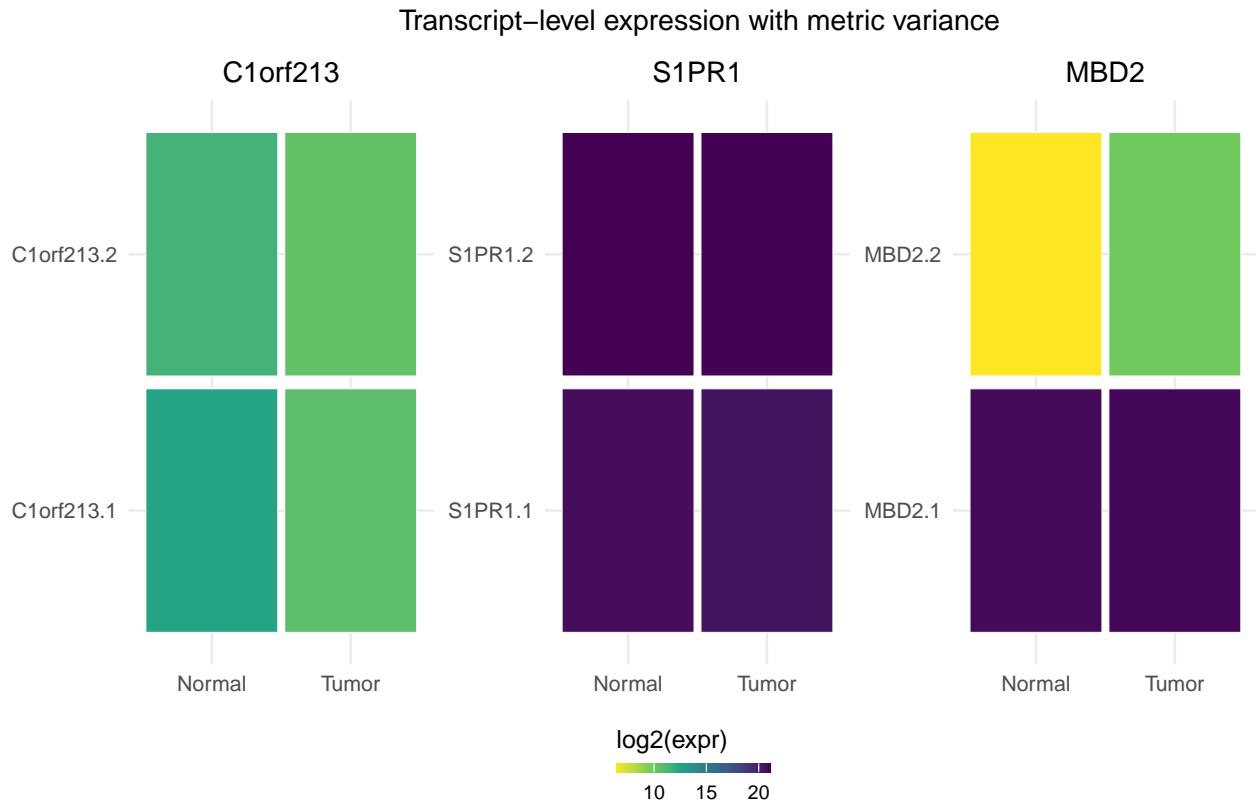
Here we will visualize the transcript-level expression for the most significant genes identified in the previous step.

```
# Plot using median aggregation; let the function pick top genes from `res`  
p_median <- plot_top_transcripts(ts_se,  
  res = res, top_n = 3, metric = "median"  
)  
print(p_median)
```



We can also plot the same top genes but using variance aggregation instead of median. This allows us to see how transcript-level expression variability differs between groups for these genes.

```
# Plot using variance aggregation
p_var <- plot_top_transcripts(ts_se,
  res = res, top_n = 3, metric = "variance"
)
print(p_var)
```



Compare between q values

Now we are going to try a different approach. Instead of evaluating the significant statistical differences by using single Tsallis q values, we will compute the normalized Tsallis entropy for two q values (0.1 and 2) to study potential scale-dependent patterns.

```
# compute Tsallis entropy for q = 1 (normalized)
q <- c(0.1, 2)
ts_se <- calculate_diversity(se,
  q = q, norm = TRUE
)
head(assay(ts_se)[1:4, 1:3])
#>      TCGA-A7-AOCH_N_q=0.1 TCGA-A7-AOCH_N_q=2 TCGA-A7-AOCH_T_q=0.1
#> MXRA8          0.8616275  0.6063802  0.6609111
#> C1orf86        0.0000000  0.0000000  0.0000000
#> PDPN          0.2657489  0.6205367  0.2588286
#> ALDH4A1        0.8111864  0.7478309  0.4667361
```

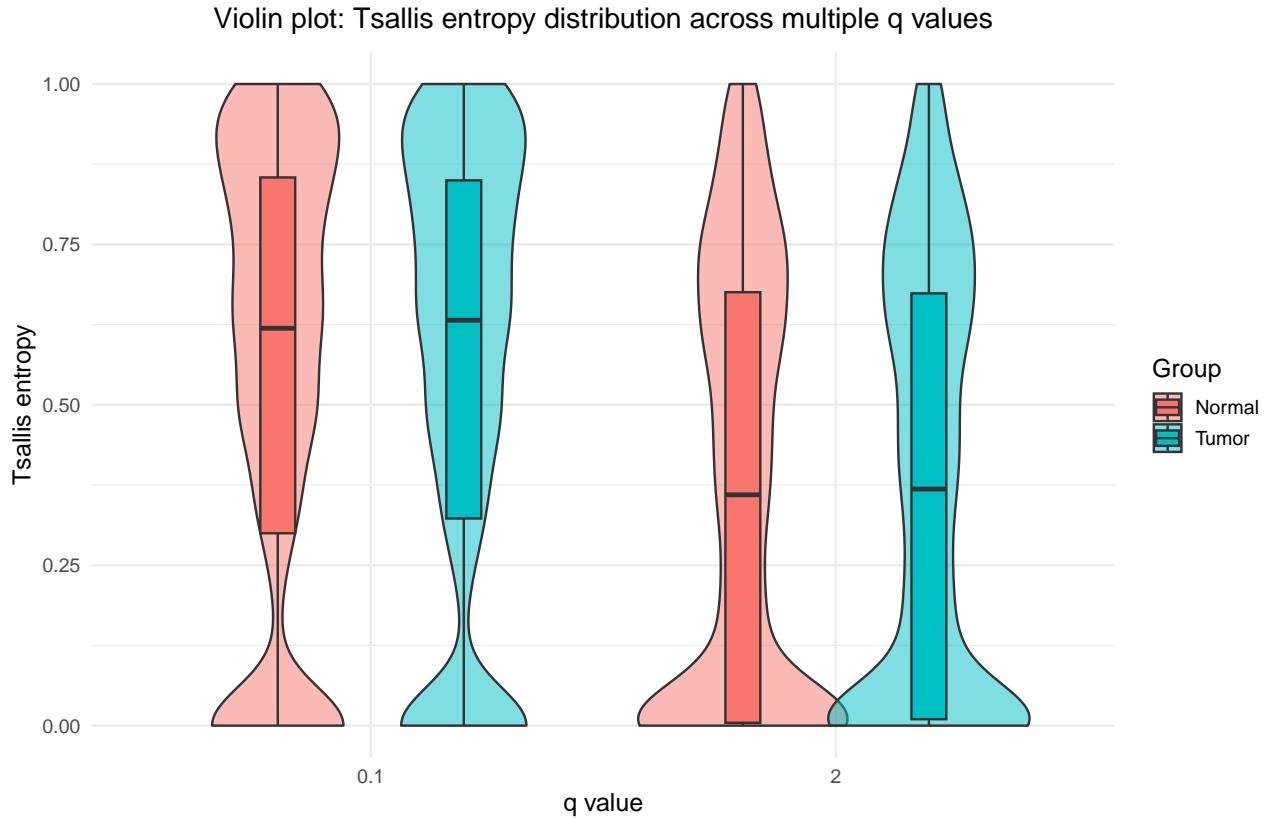
Once again, we should map the sample metadata into the multi-q `SummarizedExperiment`. This step is important to ensure that functions located downstream in the workflow can access the sample metadata correctly.

```
ts_se <- map_metadata(ts_se, coldata_df)
```

In order to show the distributional differences between groups, we are going to create **violin** and **density plots**.

The violin plot shows the distribution of Tsallis entropy values for each group and q value, allowing us to visually assess differences in diversity patterns between groups at different sensitivity levels.

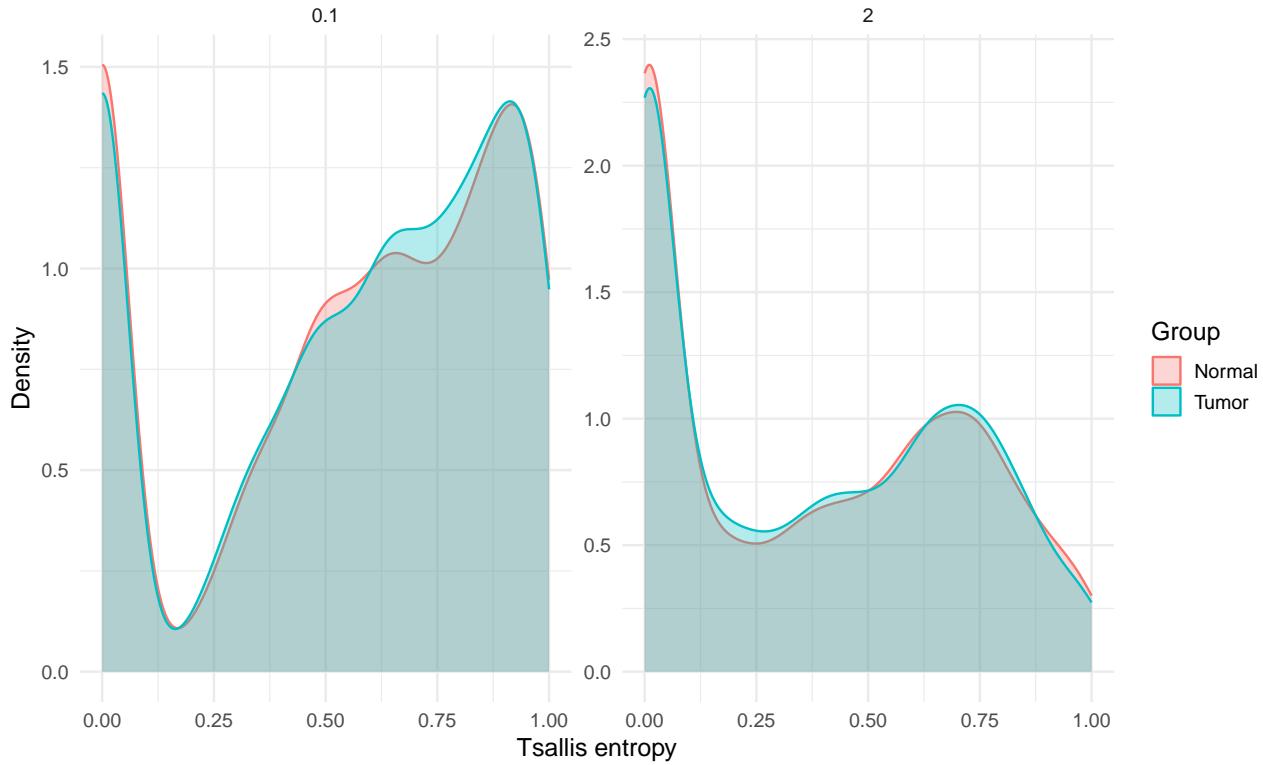
```
p_violin <- plot_tsallis_violin_multq(ts_se, assay_name = "diversity")
print(p_violin)
```



The density plot provides an equivalent view of the distribution of Tsallis entropy values, showing the density of samples across the range of diversity values for each group and q value.

```
p_density <- plot_tsallis_density_multq(ts_se, assay_name = "diversity")
print(p_density)
```

Density plot: Tsallis entropy distribution across multiple q values



Tsallis q-sequence and linear-model interaction tests

Finally, we can test for interactions between q and sample groups across a sequence of q values (**q-sequence**), that is, whether the shape of the q -curve differs between groups. This approach is a powerful way to detect scale-dependent diversity differences that may not be apparent at a single q value.

Computing Tsallis entropy across a sequence of q values is important because q acts as a sensitivity lens that shifts emphasis between rare and dominant events in data. Thus, the resulting **q-curve** can help us identify at which scales (i.e. for which q values) the diversity differences between groups are most pronounced. The Tsallis **q-sequence** can also help with quality control and outlier detection. By examining the q -curves of individual samples, we can identify samples that deviate from expected patterns, which may indicate technical issues or biological outliers.

```
# compute Tsallis entropy for a sequence of values (normalized)
qvec <- seq(0.1, 2, by = 0.05)
ts_se <- calculate_diversity(se,
  q = qvec, norm = TRUE
)
head(assay(ts_se)[1:4, 1:3])
#>           TCGA-A7-AOCH_N_q=0.1  TCGA-A7-AOCH_N_q=0.15  TCGA-A7-AOCH_N_q=0.2
#> MXRA8             0.8616275      0.8073925      0.7611249
#> C1orf86            0.0000000      0.0000000      0.0000000
#> PDPN              0.2657489      0.2738368      0.2820625
#> ALDH4A1            0.8111864      0.7506492      0.7056995
```

Map once again the sample metadata into the multi-q `SummarizedExperiment` as before.

```
# For paired designs ensure sample_base is created so LMMs can use it
ts_se <- map_metadata(ts_se, coldata_df, paired = TRUE)
```

Once we have the **q-sequence** computed, we can test for interactions between q and sample groups using one of four methods implemented in `calculate_lm_interaction()`:

- **linear**: fits a simple linear model. This is fast and interpretable and is appropriate when the q-response is approximately linear and samples are independent.
- **lmm**: fits a linear mixed model which accounts for subject-level random effects (paired or repeated measures).
- **gam**: fits a generalized additive model. This method allows to capture nonlinear q-response shapes.
- **fPCA**: uses a functional-data approach to reduce each gene's q-curve to principal component scores. This is recommended when q is densely sampled or when q-curves exhibit complex, high-dimensional variation.

Below we demonstrate a practical `lmm` example using the **Satterthwaite** approximation for per-term inference. This method is preferable for small or moderately sized paired RNA-seq studies when per-term inference is required (for example, the q:group interaction).

```
## Linear-model interaction test across q values: detect q x group interaction
if (requireNamespace("lme4", quietly = TRUE)) {
  # Compute and show top hits (by adjusted p-value)
  lm_res <- calculate_lm_interaction(ts_se,
    sample_type_col = "sample_type", method = "lmm",
    pvalue = "satterthwaite",
    subject_col = "sample_base"
  )
  # show the top entries
  print(head(lm_res, 6))
}

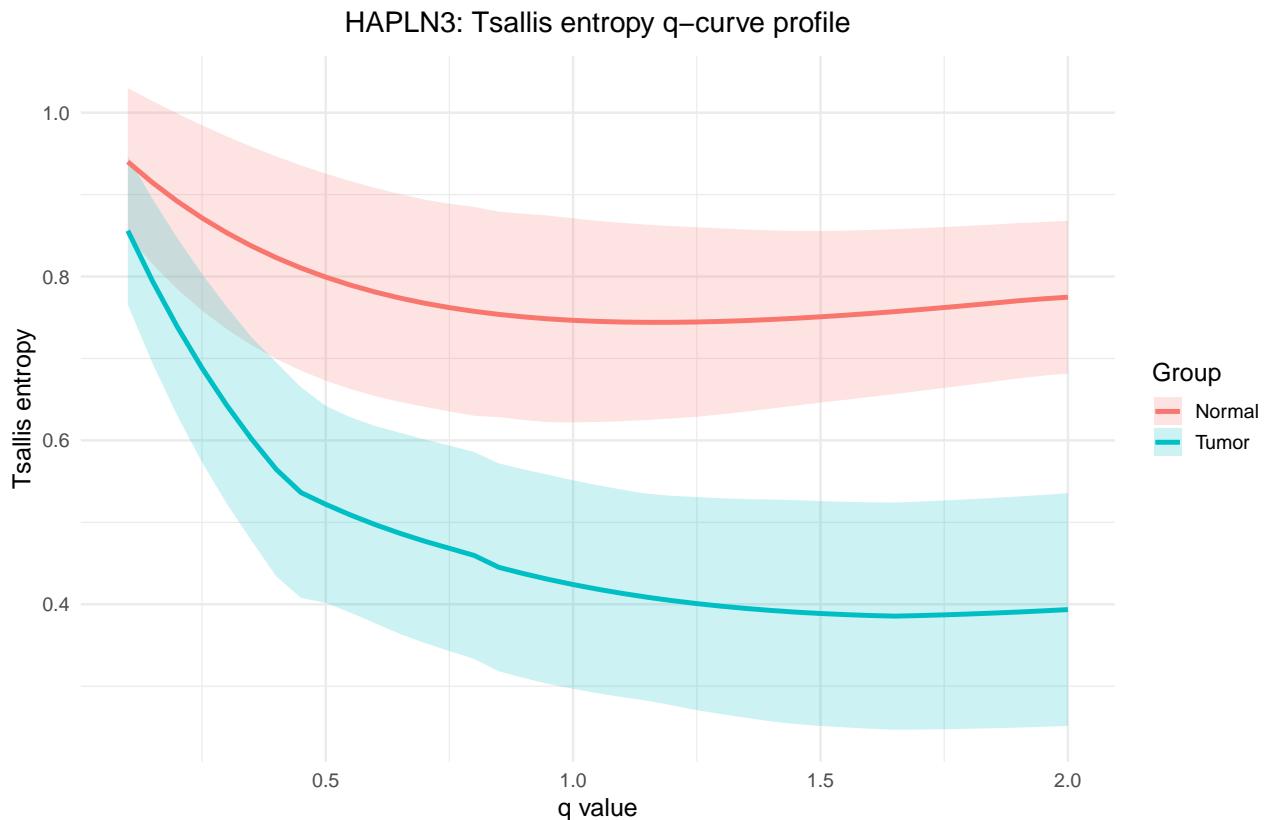
#>      gene p_interaction      p_lrt p_satterthwaite fit_method singular
#> 1  HAPLN3 2.319271e-111 2.139688e-111 2.319271e-111    lmer   FALSE
#> 2  COL1A2 8.713853e-97  8.123588e-97  8.713853e-97    lmer   FALSE
#> 3  EEF2K  3.093021e-87  2.903439e-87  3.093021e-87    lmer   FALSE
#> 4  C1orf86 8.733394e-83  8.224552e-83  8.733394e-83    lmer   FALSE
#> 5  LRRC15 7.474347e-79  7.058969e-79  7.474347e-79    lmer   FALSE
#> 6  PI16   5.493773e-73  5.210643e-73  5.493773e-73    lmer   FALSE
#>      adj_p_interaction
#> 1      4.615350e-109
#> 2      8.670284e-95
#> 3      2.051704e-85
#> 4      4.344863e-81
#> 5      2.974790e-77
#> 6      1.822101e-71
```

Tsallis q-sequence plot

With the **Tsallis q-sequence** we can produce a q-curve per sample and gene. These curves show how diversity emphasis shifts from rare to dominant isoforms as q increases and form the basis for interaction tests. The q-curve shows entropy as a function of q. Diverging curves between groups indicate **scale-dependent diversity differences**: separation at low q implies differences in rare isoforms, while separation at high q signals differences in dominant isoforms.

Now we will plot the q-curve profile for the top gene identified by the linear-model interaction test. We will start with **HAPLN3**, which codifies member of the hyaluronan and proteoglycan link protein family expressed in the extracellular matrix, and closely associated with the development and occurrence of various malignant tumors (source).

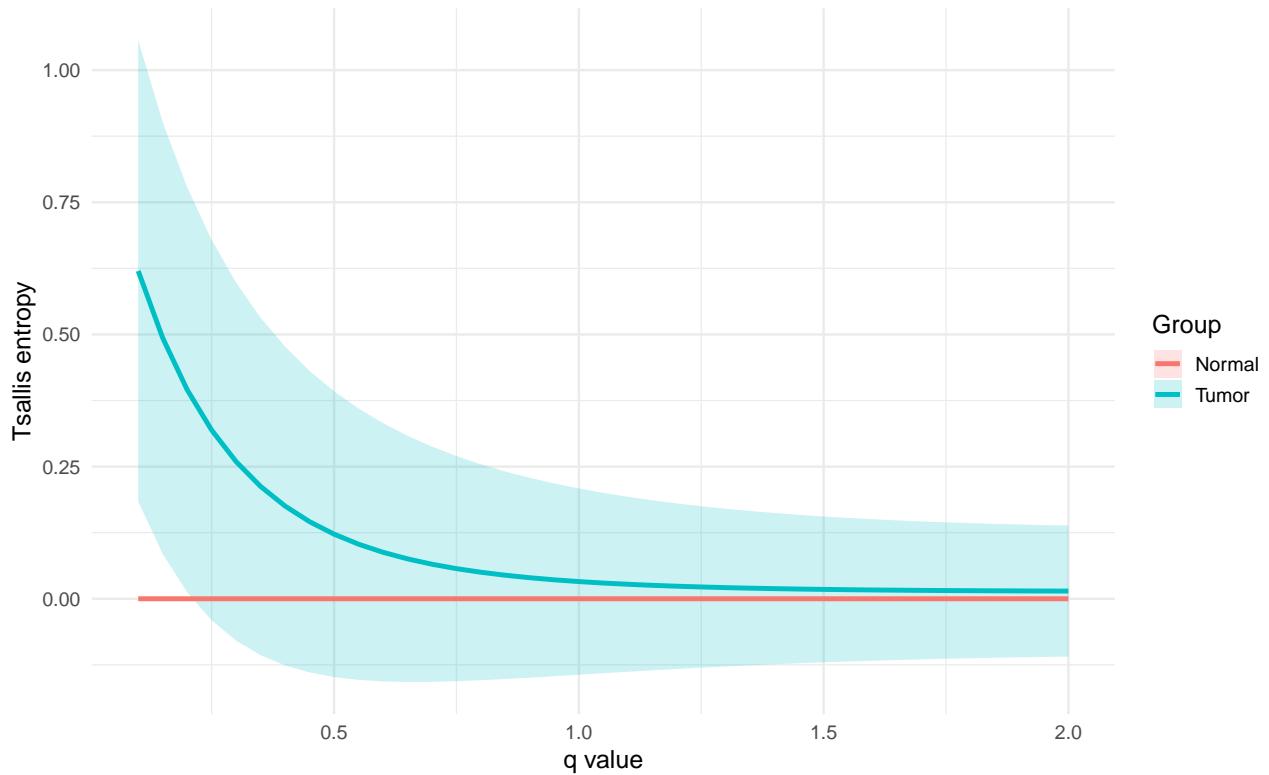
```
# Plot q-curve profile for top linear-model gene
plot_target <- plot_tsallis_gene_profile(ts_se, gene = "HAPLN3")
print(plot_target)
```



Now let's look at **COL1A2**, encodes the pro-alpha2 chain of type I collagen, a fibril-forming collagen found in most connective tissues.

```
# Plot q-curve profile for top linear-model gene
plot_target <- plot_tsallis_gene_profile(ts_se, gene = "COL1A2")
print(plot_target)
```

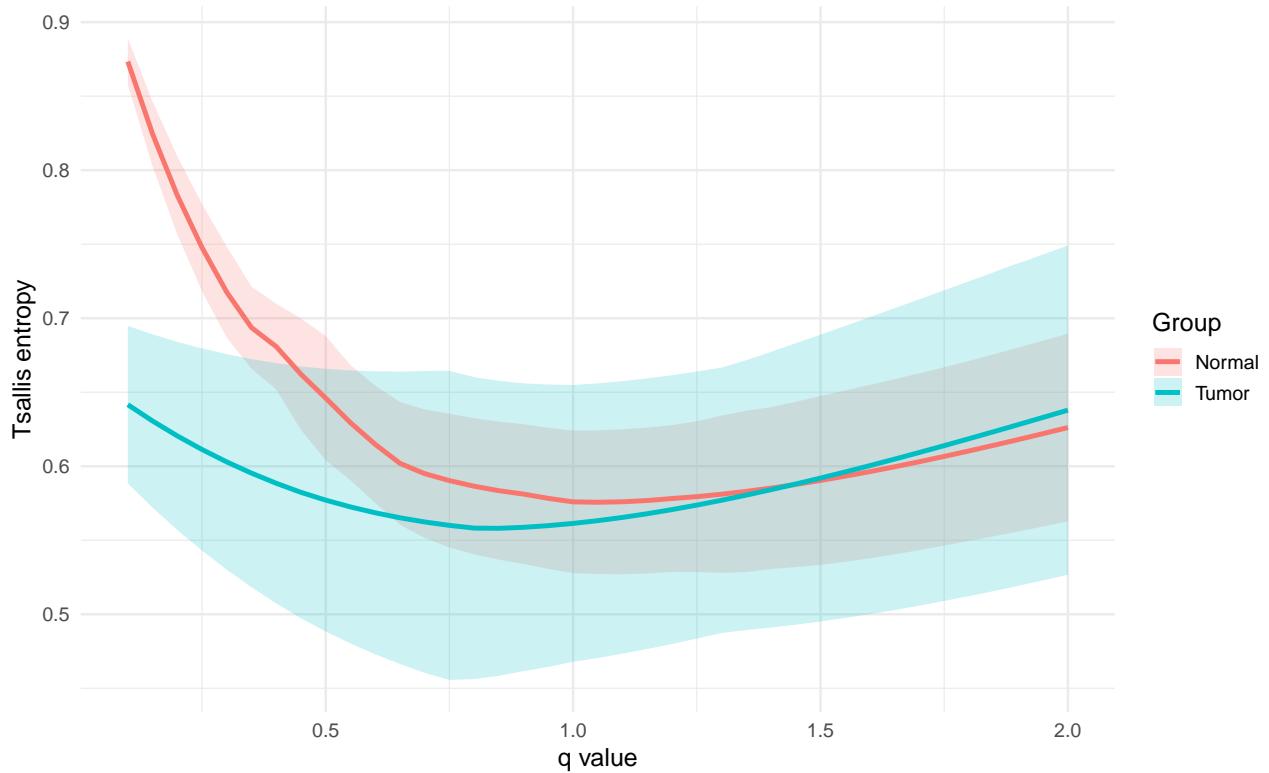
COL1A2: Tsallis entropy q–curve profile



Finally, we can visualize the gene PI16, which is emerging as an important regulator in the vascular system.

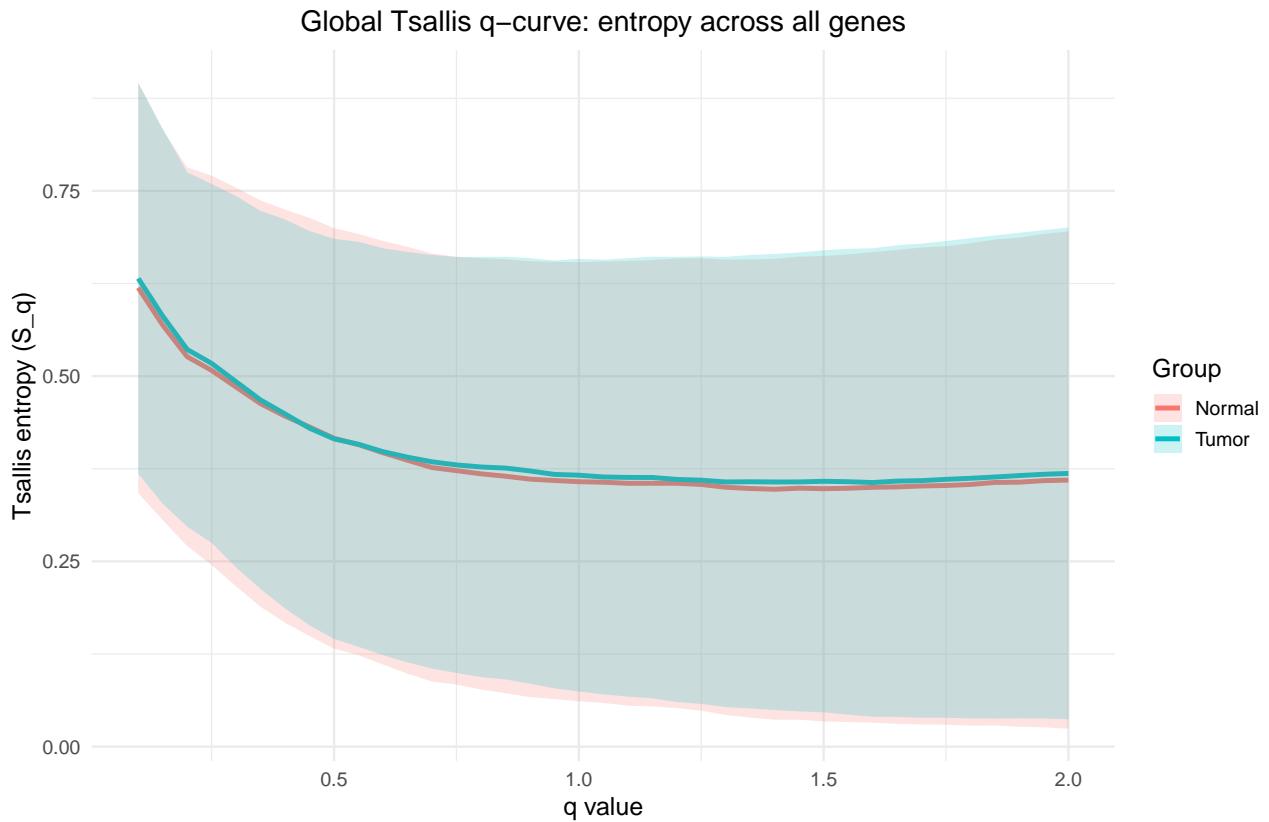
```
# Plot q-curve profile for top linear-model gene
plot_target <- plot_tsallis_gene_profile(ts_se, gene = "PI16")
print(plot_target)
```

PI16: Tsallis entropy q–curve profile



We can also plot the **overall Tsallis q-curve for all genes**. This plot provides a global view of how diversity changes across the q spectrum for the entire dataset and can reveal general trends in isoform diversity.

```
# Plot top tsallis q curve
p_qcurve <- plot_tsallis_q_curve(ts_se)
print(p_qcurve)
```



Practical notes and recommendations

Accurate isoform quantification is the foundation of reliable entropy analysis. Small counts introduce noise; missing data distorts proportions. Some important elements to keep in mind:

- **Transcript abundance filtering:** Remove very low-count transcripts to reduce technical noise (typical rule: drop transcripts with <5 reads in >5 samples, but adjust to sequencing depth).
- **Pseudocounts:** If zeros create numerical issues, you could add a small pseudocount (e.g. 1e-10). You can find more information about pseudocounts in this paper.
- **Quantification & normalization:** use bias-aware quantifiers, such Salmon or kallisto.

```
sessionInfo()
#> R version 4.5.2 (2025-10-31)
#> Platform: x86_64-conda-linux-gnu
#> Running under: Ubuntu 22.04.5 LTS
#>
#> Matrix products: default
#> BLAS/LAPACK: /home/nouser/miniconda3/lib/libopenblas-r0.3.30.so; LAPACK version 3.12.0
#>
#> locale:
#> [1] LC_CTYPE=es_ES.UTF-8          LC_NUMERIC=C
#> [3] LC_TIME=de_DE.UTF-8          LC_COLLATE=es_ES.UTF-8
#> [5] LC_MONETARY=de_DE.UTF-8      LC_MESSAGES=es_ES.UTF-8
#> [7] LC_PAPER=de_DE.UTF-8         LC_NAME=C
#> [9] LC_ADDRESS=C                LC_TELEPHONE=C
#> [11] LC_MEASUREMENT=de_DE.UTF-8 LC_IDENTIFICATION=C
```

```

#>
#> time zone: Europe/Berlin
#> tzcode source: system (glibc)
#>
#> attached base packages:
#> [1] stats4      stats       graphics   grDevices  utils      datasets   methods
#> [8] base
#>
#> other attached packages:
#> [1] SummarizedExperiment_1.40.0 Biobase_2.70.0
#> [3] GenomicRanges_1.62.1        Seqinfo_1.0.0
#> [5] IRanges_2.44.0            S4Vectors_0.48.0
#> [7] BiocGenerics_0.56.0        generics_0.1.4
#> [9] MatrixGenerics_1.22.0      matrixStats_1.5.0
#> [11] ggplot2_4.0.2            TSENAT_0.99.0
#> [13] testthat_3.3.2
#>
#> loaded via a namespace (and not attached):
#> [1] gtable_0.3.6           xfun_0.56          bslib_0.10.0
#> [4] lattice_0.22-7         numDeriv_2016.8-1.1 Rdpack_2.6.6
#> [7] vctrs_0.7.1           tools_4.5.2         tibble_3.3.1
#> [10] pkgconfig_2.0.3       Matrix_1.7-4        RColorBrewer_1.1-3
#> [13] S7_0.2.1             desc_1.4.3          lifecycle_1.0.5
#> [16] compiler_4.5.2        farver_2.1.2        brio_1.1.5
#> [19] tinytex_0.58          lmerTest_3.2-0      htmltools_0.5.9
#> [22] sass_0.4.10          yaml_2.3.12         nloptr_2.2.1
#> [25] pillar_1.11.1         jquerylib_0.1.4    tidyrr_1.3.2
#> [28] MASS_7.3-65           DelayedArray_0.36.0 cachem_1.1.0
#> [31] reformulas_0.4.4       boot_1.3-32         abind_1.4-8
#> [34] nlme_3.1-168          tidyselect_1.2.1    digest_0.6.39
#> [37] dplyr_1.2.0           purrrr_1.2.1        labeling_0.4.3
#> [40] splines_4.5.2          rprojroot_2.1.1    fastmap_1.2.0
#> [43] grid_4.5.2            cli_3.6.5           SparseArray_1.10.8
#> [46] magrittr_2.0.4          patchwork_1.3.2    S4Arrays_1.10.1
#> [49] pkgbuild_1.4.8          withr_3.0.2         scales_1.4.0
#> [52] rmarkdown_2.30          XVector_0.50.0     otel_0.2.0
#> [55] lme4_1.1-38            evaluate_1.0.5     knitr_1.51
#> [58] rbibutils_2.4.1         viridisLite_0.4.3   rlang_1.1.7
#> [61] Rcpp_1.1.1              glue_1.8.0          pkgload_1.5.0
#> [64] minqa_1.2.8            jsonlite_2.0.0     R6_2.6.1

```